

K-Nearest-Neighbors

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import confusion_matrix, accuracy_score
from sklearn.datasets import load_iris

data = load_iris()
df = pd.DataFrame(data.data, columns=data.feature_names)
df['target'] = data.target

X = df.drop('target', axis=1)
y = df['target']

print("Dataset Shape:", X.shape)
print(X.head())

Dataset Shape: (150, 4)
   sepal length (cm)   sepal width (cm)   petal length (cm)   petal width
   (cm)
0           5.1            3.5            1.4
0.2
1           4.9            3.0            1.4
0.2
2           4.7            3.2            1.3
0.2
3           4.6            3.1            1.5
0.2
4           5.0            3.6            1.4
0.2

X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)
```

- 1) KNN uses distance. If one feature is "Income" (100,000) and another is "Age" (50),
- 2) Income will dominate the distance calculation. Scaling fixes this.

```
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)
```

- 1) The "Lazy" Learner KNN is what we call a Lazy Learner. Linear/Logistic Regression are "Eager Learners." When you run .fit(), they sweat to calculate the perfect slopes (m) and intercepts (b). They "compile" the data into a mathematical equation and then throw the data away.

2) KNN does the opposite. When you run `.fit()`, it simply memorizes the data. It takes your X_{train} and y_{train} and stores them in memory.

```
k = 5
knn_model = KNeighborsClassifier(n_neighbors=k)

knn_model.fit(X_train_scaled, y_train)
KNeighborsClassifier()
```

- 1) While there is no mathematical optimization (like Gradient Descent), there is a computer science optimization happening in the background of `.fit()`.
- 2) If you have 100,000 points, calculating the distance to every single one later would be too slow. So, sklearn organizes the data into a smart tree structure, usually a K-D Tree or Ball Tree.
- 3) Think of this like organizing a library. Instead of throwing books in a pile, `.fit()` organizes them by genre, then author, then title.
- 4) This makes finding the "nearest neighbors" much faster later on.

```
y_pred = knn_model.predict(X_test_scaled)

acc = accuracy_score(y_test, y_pred)
print(f"Accuracy: {acc:.2%}")


Accuracy: 100.00%
```

- 1) Since `.fit()` was lazy, the `.predict()` function has to do all the heavy lifting. When you ask it to predict a flower species:
Distance Calculation: It calculates the Euclidean Distance between your new flower (x_{new}) and the stored training points (x_i).

$$d(x, y) = \sqrt{\sum_{i=1}^n (x_i - y_i)^2}$$

- 2) It's just the Pythagorean theorem ($a^2 + b^2 = c^2$) extended to 4 dimensions (since we have 4 features).
Sorting: It finds the K points with the smallest distance d .
Voting: It counts the classes of those K neighbors and returns the winner.

```
plt.figure(figsize=(6, 4))
sns.heatmap(confusion_matrix(y_test, y_pred), annot=True, fmt='d',
cmap='Blues')
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title('Confusion Matrix')
plt.show()
```




- 1) Low K (e.g., K=1): The model is too complex. It memorizes noise. This is Overfitting (High Variance).
- 2) High K (e.g., K=50): The model is too simple. It ignores details and just votes for the majority class. This is Underfitting (High Bias).
- 3) The "Valley" (Optimal K): The point in the middle where you perfectly balance bias and variance.

```
# Finding the optimal K
# Hyperparameter Tuning
error_rate = []

for i in range(1, 20):
    knn = KNeighborsClassifier(n_neighbors=i)
    knn.fit(X_train_scaled, y_train)
    pred_i = knn.predict(X_test_scaled)
    # Average of where prediction was NOT equal to actual
    error_rate.append(np.mean(pred_i != y_test))

# Plotting
plt.figure(figsize=(10,6))
plt.plot(range(1, 20), error_rate, color='blue', linestyle='dashed',
marker='o')
plt.title('Error Rate vs. K Value')
plt.xlabel('K')
```

```
plt.ylabel('Error Rate')
plt.show()
```



- 1) The iris dataset is well classified dataset so for any value of K here doesn't affect the accuracy.
- 2) But for any another classification dataset for KNN you must choose the optimal K value for which error rate is minimum.
- 3) On paper this value of K can be obtained by K-fold cross validation technique using random search.